

EVALUATING WINE QUALITY VIA PHYSICOCHEMICAL TESTS

MATH 2319 Machine Learning Applied Project Phase II

HUYNH AI LOAN (s3655461)

11 June 2018

- 1 Introduction
- 2 Methodology
- 3 Hyperparameter Tune-Fining
 - 3.1 K-Nearest Neighbour
 - 3.2 Support Vector Machine
 - 3.3 Random Forest
 - 3.4 Threshold Adjustment
 - 3.4.1 KNN
 - 3.4.2 Suport Vector Machine
 - 3.4.3 Random Forest
- 4 Evaluation
- 5 Discussion
- 6 Conclusion
- 7 References

1 Introduction

The objective of this project is to build classifiers to predict whether physicochemical tests make the quality of wine larger than 5 grade in range of score between 0 (very bad) and 10 (very excellent) which are made by wine experts. The data sets were collected from the UCI Machine Learning Repository. In Phase I, we cleaned the data and re-categorised some descriptive features to be less granular. In Phase II, we built three binary-classifiers on the cleaned data. Section 2 describes an overview of our methodology. Section 3 discusses the classifiers and their tuning process. Section 4 compares the performance of the classifiers using the same resampling method. The last section concludes with a summary.

2 Methodology

In this report, the three classifiers - Random Forest (RF), K-Nearest Neighbour (KNN) and Support Vector Machine (SVM) are considered to deal with the problem. The target feature in dataset was grouped into two levels which were less than or equal 5 (≤ 5) and larger than 5 (> 5). The dataset was splitted into training and test set with ratio 7:3. Each classifier was trained to make probability predictions in order that we could adjust prediction threshold to evaluate the performance. For fine-tuning process, we used 5-folded cross validation stratified sampling on each classifier.

Using the tuned hyperparameters and the optimal thresholds defined from previous steps, we made prediction on the test data for each classifier. We used mean misclassification error rate (mmce) and confusion matrix on the test data to evaluate the classifiers's performance.

3 Hyperparameter Tune-Fining

3.1 K-Nearest Neighbour

KNN uses distance to classify the features. Therefore, it is necessary to standardize the predictor variables. There are two type of distances used in this report including Manhattan and Euclidian distance. In addition, we also ran a grid search on k values in range from 0 to 10 in order to define which the best k neighbours give the best result for model. The tuning result is as below:

[Hide](#)

```
inTrain <- createDataPartition(cleaned_data$quality, p = 0.7, list = FALSE)
training_data <- cleaned_data[inTrain,]
test_data <- cleaned_data[-inTrain,]
task <- makeClassifTask(data = training_data, target = 'quality', id = 'wine')
knn_learner <- makeLearner('classif.kknn', predict.type = 'prob')
ps_knn <- makeParamSet(
  makeDiscreteParam('k', values = seq(2, 10, by = 1)),
  makeDiscreteParam('distance', values=c(1,2)),
  makeDiscreteParam('kernel', values = "cos")
)
ctrl <- makeTuneControlGrid()
rdesc <- makeResampleDesc("CV", iters = 5L, stratify = TRUE)
# Configure tune wrapper with tune-tuning settings
knn_tunedLearner <- makeTuneWrapper(learner = knn_learner, resampling = rdesc, measures = mmce, par.set= ps_knn, control = ctrl)
# # Train the tune wrappers
knn_tuneWrapper <- mlr::train(knn_tunedLearner, task)
```

```
[Tune] Started tuning learner classif.kknn for parameter set:
      Type len Def          Constr Req Tunable Trafo
k      discrete - - 2,3,4,5,6,7,8,9,10 - TRUE -
distance discrete - -          1,2 - TRUE -
kernel  discrete - -          cos - TRUE -
With control class: TuneControlGrid
Imputation value: 1
[Tune-x] 1: k=2; distance=1; kernel=cos
[Tune-y] 1: mmce.test.mean=0.2192030; time: 0.0 min
[Tune-x] 2: k=3; distance=1; kernel=cos
[Tune-y] 2: mmce.test.mean=0.2157200; time: 0.0 min
[Tune-x] 3: k=4; distance=1; kernel=cos
[Tune-y] 3: mmce.test.mean=0.2154045; time: 0.0 min
[Tune-x] 4: k=5; distance=1; kernel=cos
[Tune-y] 4: mmce.test.mean=0.2103377; time: 0.0 min
[Tune-x] 5: k=6; distance=1; kernel=cos
[Tune-y] 5: mmce.test.mean=0.2112876; time: 0.0 min
[Tune-x] 6: k=7; distance=1; kernel=cos
[Tune-y] 6: mmce.test.mean=0.2103362; time: 0.0 min
[Tune-x] 7: k=8; distance=1; kernel=cos
[Tune-y] 7: mmce.test.mean=0.2100218; time: 0.0 min
[Tune-x] 8: k=9; distance=1; kernel=cos
[Tune-y] 8: mmce.test.mean=0.2093884; time: 0.0 min
[Tune-x] 9: k=10; distance=1; kernel=cos
[Tune-y] 9: mmce.test.mean=0.2084405; time: 0.0 min
[Tune-x] 10: k=2; distance=2; kernel=cos
[Tune-y] 10: mmce.test.mean=0.2220536; time: 0.0 min
[Tune-x] 11: k=3; distance=2; kernel=cos
[Tune-y] 11: mmce.test.mean=0.2195235; time: 0.0 min
[Tune-x] 12: k=4; distance=2; kernel=cos
[Tune-y] 12: mmce.test.mean=0.2163534; time: 0.0 min
[Tune-x] 13: k=5; distance=2; kernel=cos
[Tune-y] 13: mmce.test.mean=0.2128704; time: 0.0 min
[Tune-x] 14: k=6; distance=2; kernel=cos
[Tune-y] 14: mmce.test.mean=0.2100228; time: 0.0 min
[Tune-x] 15: k=7; distance=2; kernel=cos
[Tune-y] 15: mmce.test.mean=0.2090724; time: 0.0 min
[Tune-x] 16: k=8; distance=2; kernel=cos
[Tune-y] 16: mmce.test.mean=0.2062208; time: 0.0 min
[Tune-x] 17: k=9; distance=2; kernel=cos
[Tune-y] 17: mmce.test.mean=0.2090749; time: 0.0 min
[Tune-x] 18: k=10; distance=2; kernel=cos
[Tune-y] 18: mmce.test.mean=0.2065433; time: 0.0 min
[Tune] Result: k=8; distance=2; kernel=cos : mmce.test.mean=0.2062208
```

[Hide](#)

```
# Get Tune Result
print(getTuneResult(knn_tuneWrapper))
```

```
Tune result:
Op. pars: k=8; distance=2; kernel=cos
mmce.test.mean=0.2062208
```

3.2 Support Vector Machine

We considered `gamma` and `cost` parameters for tuning. The `gamma` parameter defines how far the influence of a training data reaches. The higher value of `gamma` will try to fit training dataset. The `cost` of constraint violation controls the trade off between smooth decision boundary and classifying the training points correctly. We experimented with `gamma` value from 0.5 to 3 and `cost` value in range of (0.5, 3).

[Hide](#)

```
svm_learner <- makeLearner('classif.svm', predict.type = 'prob')
ps_svm <- makeParamSet(
  makeDiscreteParam('gamma', values = c(0.5,1, 1.5, 2, 2.5, 3)),
  makeDiscreteParam('cost', values = c(0.5,1, 1.5, 2, 2.5, 3))
)
svm_tunedLearner <- makeTuneWrapper(svm_learner, rdesc, measures=list(acc,mmce), ps_svm, ctrl
1)
svm_tuneWrapper <- mlr::train(svm_tunedLearner, task)
```

[Tune] Started tuning learner `classif.svm` for parameter set:

	Type	len	Def	Constr	Req	Tunable	Trafo
<code>gamma</code>	discrete	-	-	0.5,1,1.5,2,2.5,3	-	TRUE	-
<code>cost</code>	discrete	-	-	0.5,1,1.5,2,2.5,3	-	TRUE	-

With control class: `TuneControlGrid`

Imputation value: -0Imputation value: 1

[Tune-x] 1: `gamma=0.5; cost=0.5`

[Tune-y] 1: `acc.test.mean=0.7963174,mmce.test.mean=0.2036826; time: 0.3 min`

[Tune-x] 2: `gamma=1; cost=0.5`

[Tune-y] 2: `acc.test.mean=0.7947371,mmce.test.mean=0.2052629; time: 0.4 min`

[Tune-x] 3: `gamma=1.5; cost=0.5`

[Tune-y] 3: `acc.test.mean=0.7836541,mmce.test.mean=0.2163459; time: 0.4 min`

[Tune-x] 4: `gamma=2; cost=0.5`

[Tune-y] 4: `acc.test.mean=0.7773205,mmce.test.mean=0.2226795; time: 0.4 min`

[Tune-x] 5: `gamma=2.5; cost=0.5`

[Tune-y] 5: `acc.test.mean=0.7738355,mmce.test.mean=0.2261645; time: 0.4 min`

[Tune-x] 6: `gamma=3; cost=0.5`

[Tune-y] 6: `acc.test.mean=0.7754203,mmce.test.mean=0.2245797; time: 0.4 min`

[Tune-x] 7: `gamma=0.5; cost=1`

[Tune-y] 7: `acc.test.mean=0.8020186,mmce.test.mean=0.1979814; time: 0.4 min`

[Tune-x] 8: `gamma=1; cost=1`

[Tune-y] 8: `acc.test.mean=0.7944201,mmce.test.mean=0.2055799; time: 0.5 min`

[Tune-x] 9: `gamma=1.5; cost=1`

[Tune-y] 9: `acc.test.mean=0.7817549,mmce.test.mean=0.2182451; time: 0.5 min`

[Tune-x] 10: `gamma=2; cost=1`

[Tune-y] 10: `acc.test.mean=0.7773215,mmce.test.mean=0.2226785; time: 0.5 min`

[Tune-x] 11: `gamma=2.5; cost=1`

[Tune-y] 11: `acc.test.mean=0.7744694,mmce.test.mean=0.2255306; time: 0.5 min`

[Tune-x] 12: `gamma=3; cost=1`

[Tune-y] 12: `acc.test.mean=0.7747869,mmce.test.mean=0.2252131; time: 0.5 min`

[Tune-x] 13: `gamma=0.5; cost=1.5`

[Tune-y] 13: `acc.test.mean=0.8048712,mmce.test.mean=0.1951288; time: 0.4 min`

[Tune-x] 14: `gamma=1; cost=1.5`

[Tune-y] 14: `acc.test.mean=0.7966383,mmce.test.mean=0.2033617; time: 0.5 min`

[Tune-x] 15: `gamma=1.5; cost=1.5`

3.3 Random Forest

For RF, we did experiment with `mtry` of 1 through 10. The result is as below:

Hide

```
rf_learner <- makeLearner('classif.randomForest', predict.type = 'prob')
ps_rf <- makeParamSet(
  makeDiscreteParam('mtry', values = seq(1,10, by = 1))
)
rf_tunedLearner <- makeTuneWrapper(rf_learner, rdesc, measures=list(acc,mmce), ps_rf, ctrl)
rf_tuneWrapper <- mlr::train(rf_tunedLearner, task)
```

```
[Tune] Started tuning learner classif.randomForest for parameter set:
      Type len Def          Constr Req Tunable Trafo
mtry discrete - - 1,2,3,4,5,6,7,8,9,10 - TRUE -
With control class: TuneControlGrid
Imputation value: -0Imputation value: 1
[Tune-x] 1: mtry=1
[Tune-y] 1: acc.test.mean=0.8305422,mmce.test.mean=0.1694578; time: 0.1 min
[Tune-x] 2: mtry=2
[Tune-y] 2: acc.test.mean=0.8359285,mmce.test.mean=0.1640715; time: 0.1 min
[Tune-x] 3: mtry=3
[Tune-y] 3: acc.test.mean=0.8321265,mmce.test.mean=0.1678735; time: 0.2 min
[Tune-x] 4: mtry=4
[Tune-y] 4: acc.test.mean=0.8305422,mmce.test.mean=0.1694578; time: 0.2 min
[Tune-x] 5: mtry=5
[Tune-y] 5: acc.test.mean=0.8286405,mmce.test.mean=0.1713595; time: 0.2 min
[Tune-x] 6: mtry=6
[Tune-y] 6: acc.test.mean=0.8235727,mmce.test.mean=0.1764273; time: 0.2 min
[Tune-x] 7: mtry=7
[Tune-y] 7: acc.test.mean=0.8232572,mmce.test.mean=0.1767428; time: 0.2 min
[Tune-x] 8: mtry=8
[Tune-y] 8: acc.test.mean=0.8232582,mmce.test.mean=0.1767418; time: 0.2 min
[Tune-x] 9: mtry=9
[Tune-y] 9: acc.test.mean=0.8251585,mmce.test.mean=0.1748415; time: 0.2 min
[Tune-x] 10: mtry=10
[Tune-y] 10: acc.test.mean=0.8194563,mmce.test.mean=0.1805437; time: 0.2 min
[Tune] Result: mtry=2 : acc.test.mean=0.8359285,mmce.test.mean=0.1640715
```

Hide

```
# Get Tune Result
print(getTuneResult(rf_tuneWrapper))
```

```
Tune result:
Op. pars: mtry=2
acc.test.mean=0.8359285,mmce.test.mean=0.1640715
```

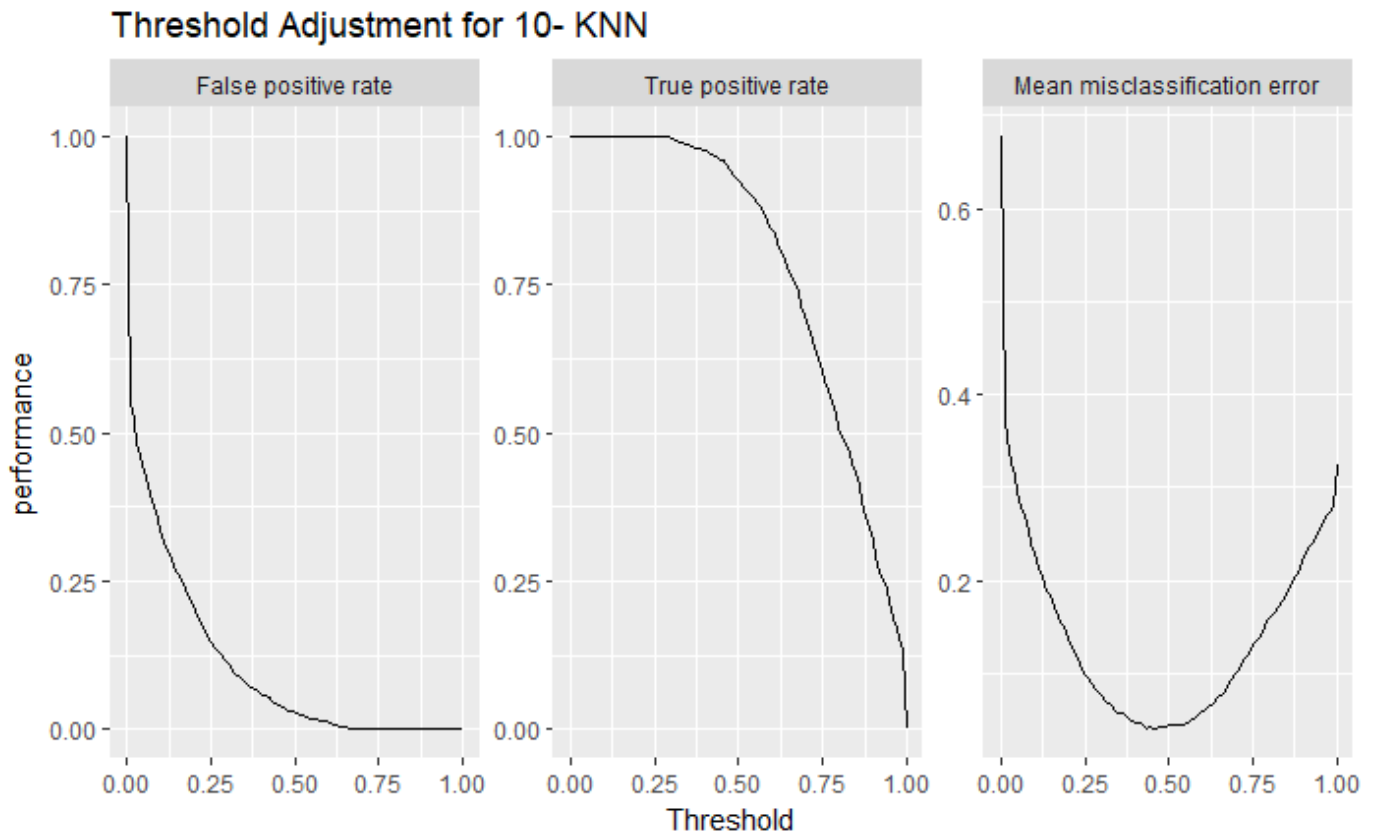
3.4 Threshold Adjustment

The following figures show the value of mmce vs the range of probability thresholds. The thresholds which may be used to determine the probability of wine with above average quality (quality >5) were approximately 0.45, 0.28 and 0.37 for 10-KNN, SVM and RF respectively.

3.4.1 KNN

Hide

```
# Predict on training data
knn_tunePredict <- predict(knn_tuneWrapper, task)
# Get threshold values for KNN learner ----
dt_knn_thresholds <- generateThreshVsPerfData(knn_tunePredict, measures = list(fpr, tpr, mmce))
# Plot thresholds adjustment for each learner
mlr::plotThreshVsPerf(dt_knn_thresholds) + labs(title = 'Threshold Adjustment for 10- KNN', x = 'Threshold')
```

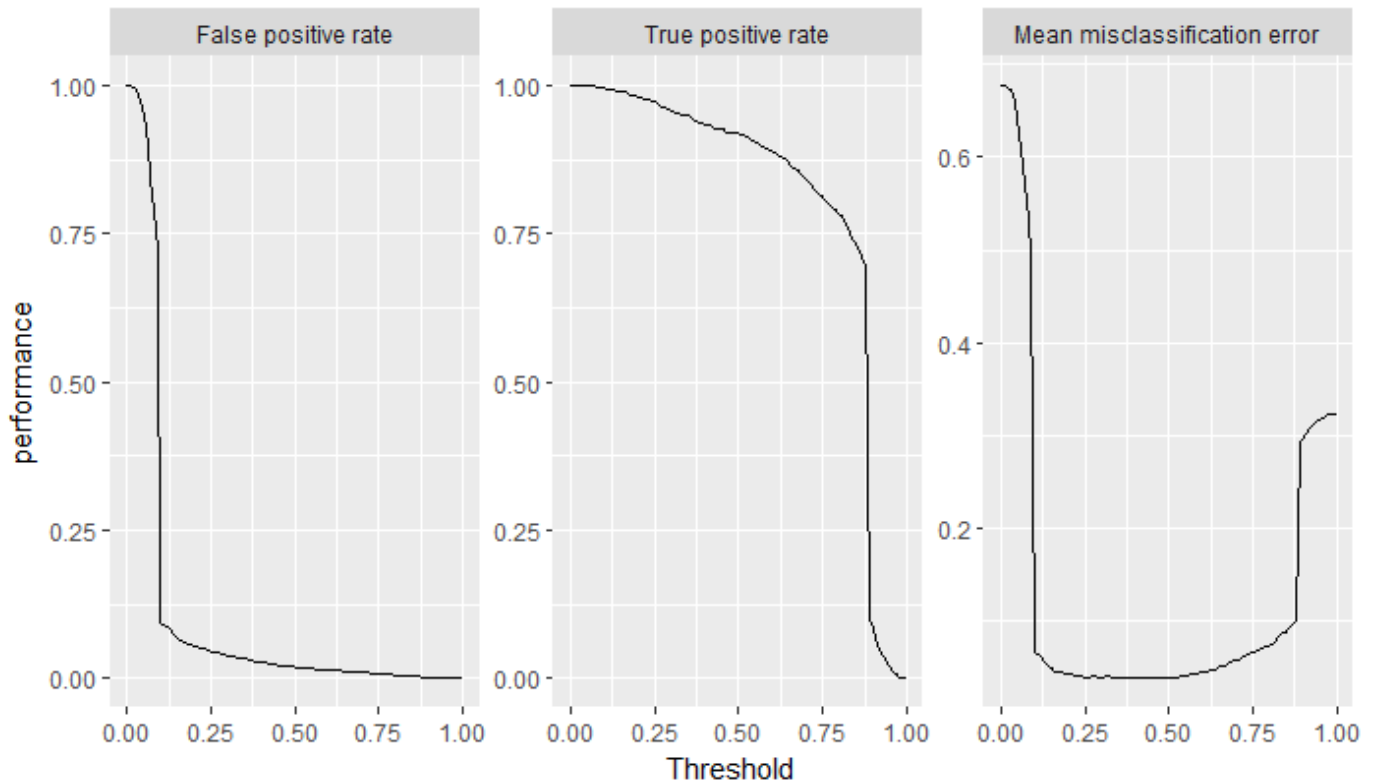

[Hide](#)

```
# Get Threshold value of KNN
knn_threshold <- dt_knn_thresholds$data$threshold[ which.min(dt_knn_thresholds$data$mmce) ]
knn_threshold
```

```
[1] 0.4545455
```

3.4.2 Support Vector Machine

Threshold Adjustment for Support Vector Machine


[Hide](#)

```
# Get Threshold value of SVM:
svm_threshold<- dt_svm_thresholds$data$threshold[ which.min(dt_svm_thresholds$data$mmce) ]
svm_threshold
```

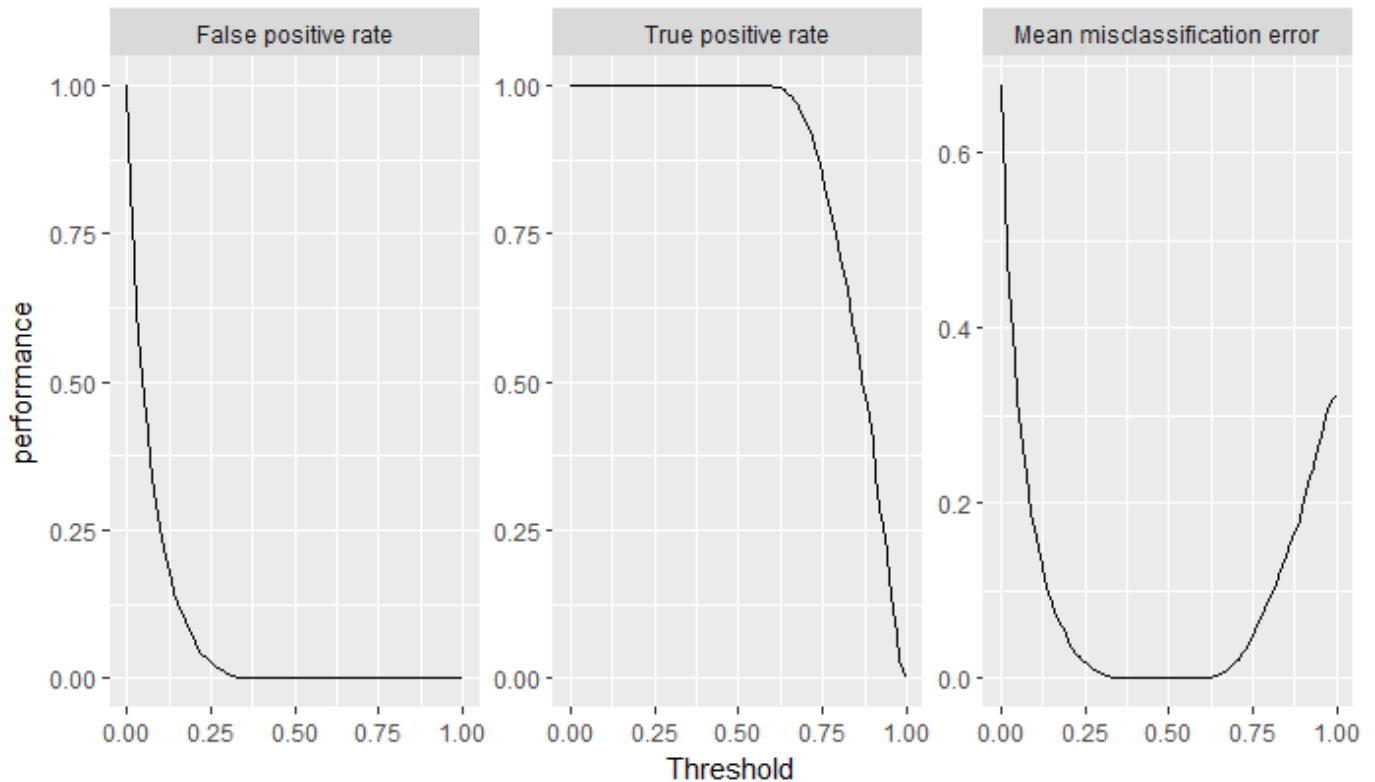
```
[1] 0.3535354
```

3.4.3 Random Forest

[Hide](#)

```
rf_tunePredict <- predict(rf_tuneWrapper, task)
dt_rf_thresholds <- generateThreshVsPerfData(rf_tunePredict, measures = list(fpr, tpr, mmce))
mlr::plotThreshVsPerf(dt_rf_thresholds) + labs(title = 'Threshold Adjustment for Random Forest', x = 'Threshold')
```

Threshold Adjustment for Random Forest

[Hide](#)

```
# Get threshold for RF
rf_threshold<- dt_rf_thresholds$data$threshold[ which.min(dt_rf_thresholds$data$mmce) ]
rf_threshold
```

```
[1] 0.3636364
```

4 Evaluation

Making prediction on test data for each classifier


```
[1] "==== Predict test data using KNN ====="
Prediction: 1352 observations
predict.type: prob
threshold: <=5=0.45,>=5=0.55
time: 0.20
  truth  prob.<=5  prob.>=5  response
1    >=5  0.08908340  0.9109166    >=5
2    >=5  0.03434164  0.9656584    >=5
12   >=5  0.29859661  0.7014034    >=5
14   >=5  0.00000000  1.0000000    >=5
16   >=5  0.25270199  0.7472980    >=5
17   <=5  0.65160903  0.3483910    <=5
... (#rows: 1352, #cols: 4)
[1] "==== Predict test data using SVM ====="
Prediction: 1352 observations
predict.type: prob
threshold: <=5=0.35,>=5=0.65
time: 0.30
  truth  prob.>=5  prob.<=5  response
1    >=5  0.9037137  0.09628633    >=5
2    >=5  0.9036612  0.09633884    >=5
12   >=5  0.6489929  0.35100713    >=5
14   >=5  0.9476377  0.05236232    >=5
16   >=5  0.7202245  0.27977545    >=5
17   <=5  0.1300488  0.86995116    <=5
... (#rows: 1352, #cols: 4)
[1] "==== Predict test data using RF ====="
Prediction: 1352 observations
predict.type: prob
threshold: <=5=0.36,>=5=0.64
time: 0.09
  truth  prob.<=5  prob.>=5  response
1    >=5    0.272    0.728    >=5
2    >=5    0.174    0.826    >=5
12   >=5    0.194    0.806    >=5
14   >=5    0.090    0.910    >=5
16   >=5    0.410    0.590    <=5
17   <=5    0.700    0.300    <=5
... (#rows: 1352, #cols: 4)
```

Using the parameters and threshold levels, we calculated the ROC measures for each classifier. The Confusion Matrix and ROC measures of KNN classifier is as follow:

```
[1] "=====  
Confusion Matrix=====
```

Relative confusion matrix (normalized by row/column):

	predicted		
true	<=5	>=5	-err.-
<=5	0.64/0.67	0.36/0.17	0.36
>=5	0.15/0.33	0.85/0.83	0.15
-err.-	0.33	0.17	0.22

Absolute confusion matrix:

	predicted		
true	<=5	>=5	-err.-
<=5	278	159	159
>=5	140	775	140
-err.-	140	159	299

```
[1] "=====  
ROC Measures =====
```

	predicted		
true	<=5	>=5	
<=5	278	159	tpr: 0.64 fnr: 0.36
>=5	140	775	fpr: 0.15 tnr: 0.85
			ppv: 0.67 for: 0.17 lrp: 4.16 acc: 0.78
			fdr: 0.33 npv: 0.83 lrm: 0.43 dor: 9.68

Abbreviations:

tpr - True positive rate (Sensitivity, Recall)
 fpr - False positive rate (Fall-out)
 fnr - False negative rate (Miss rate)
 tnr - True negative rate (Specificity)
 ppv - Positive predictive value (Precision)
 for - False omission rate
 lrp - Positive likelihood ratio (LR+)
 fdr - False discovery rate
 npv - Negative predictive value
 acc - Accuracy
 lrm - Negative likelihood ratio (LR-)
 dor - Diagnostic odds ratio

The Confusion Matrix and ROC measures of SVM classifier is as follow:

```
[1] "=====  
Confusion Matrix=====
```

Relative confusion matrix (normalized by row/column):

	predicted		
true	<=5	>=5	-err.-
<=5	0.62/0.74	0.38/0.17	0.38
>=5	0.10/0.26	0.90/0.83	0.10
-err.-	0.26	0.17	0.19

Absolute confusion matrix:

	predicted		
true	<=5	>=5	-err.-
<=5	270	167	167
>=5	96	819	96
-err.-	96	167	263

```
[1] "=====  
ROC Measures =====
```

	predicted		
true	<=5	>=5	
<=5	270	167	tpr: 0.62 fnr: 0.38
>=5	96	819	fpr: 0.1 tnr: 0.9
			ppv: 0.74 for: 0.17 lrp: 5.89 acc: 0.81
			fdr: 0.26 npv: 0.83 lrm: 0.43 dor: 13.79

Abbreviations:

tpr - True positive rate (Sensitivity, Recall)
 fpr - False positive rate (Fall-out)
 fnr - False negative rate (Miss rate)
 tnr - True negative rate (Specificity)
 ppv - Positive predictive value (Precision)
 for - False omission rate
 lrp - Positive likelihood ratio (LR+)
 fdr - False discovery rate
 npv - Negative predictive value
 acc - Accuracy
 lrm - Negative likelihood ratio (LR-)
 dor - Diagnostic odds ratio

The Confusion Matrix and ROC measures of RF classifier is as follow:

```
[1] "=====  
Confusion Matrix=====
```

Relative confusion matrix (normalized by row/column):

	predicted		
true	<=5	>=5	-err.-
<=5	0.63/0.77	0.37/0.16	0.37
>=5	0.09/0.23	0.91/0.84	0.09
-err.-	0.23	0.16	0.18

Absolute confusion matrix:

	predicted		
true	<=5	>=5	-err.-
<=5	277	160	160
>=5	81	834	81
-err.-	81	160	241

```
[1] "=====  
ROC Measures =====
```

	predicted		
true	<=5	>=5	
<=5	277	160	tpr: 0.63 fnr: 0.37
>=5	81	834	fpr: 0.09 tnr: 0.91
			ppv: 0.77 for: 0.16 lrp: 7.16 acc: 0.82
			fdr: 0.23 npv: 0.84 lrm: 0.4 dor: 17.83

Abbreviations:

tpr - True positive rate (Sensitivity, Recall)
 fpr - False positive rate (Fall-out)
 fnr - False negative rate (Miss rate)
 tnr - True negative rate (Specificity)
 ppv - Positive predictive value (Precision)
 for - False omission rate
 lrp - Positive likelihood ratio (LR+)
 fdr - False discovery rate
 npv - Negative predictive value
 acc - Accuracy
 lrm - Negative likelihood ratio (LR-)
 dor - Diagnostic odds ratio

It is obviously to see that RandomForest classifier gave higher accuracy rate rather than KNN and SVM.

5 Discussion

Three models gave the accuracy more than 79%. But RandomForest produced the better performance. All three classifiers did perform high accuracy in predicting the quality of wine larger than 5. However, three models cannot deal with imbalance issues in this dataset.

In addition, the time execution for SVM is longer than KNN and RF. It might imply that SVM may be not suitable for large dataset.

6 Conclusion

Among three classifiers, the Random Forest produces the best performance in predicting whether physicochemical tests give the quality of wine larger than 5 grade of score between 0 (very bad) and 10 (very excellent) which were evaluated by wine experts. We split the dataset into training and test sets with ratio 7:3.

Based on that, we determined the optimal value of the selected hyperparameters of each classifier and the probability threshold. In the future work, we will consider another solution to deal with imbalance issues from this dataset.

7 References